

Java NIO (New I/O) and Non-Blocking Sockets

New I/O ([java.nio.*](#))

- **New I/O APIs introduced in JDK 1.4**
- **Provides a new I/O model based on channels, buffers and selectors**
- **Enables non-blocking I/O**
- **Allows improving performance of distributed applications (mostly for the server side)**

Features in NIO APIs

- *Buffers* for data of primitive types, e.g. char, int
- *Channels*, a new primitive I/O abstraction
- *A multiplexed, non-blocking I/O facility*
(selectors, selection keys, selectable channels)
for writing scalable servers
- *Character-set encoders and decoders*
- *A pattern-matching facility* based on Perl-style
regular expressions (`java.util`)
- A file interface that supports locks and memory
mapping

NIO Programming Abstractions


- *Buffers*
 - Containers for data
 - Can be filled, drained, flipped, rewind, etc.
 - Can be written/read to/from a channel
- *Channels* of various types
 - Represent connections to entities capable of performing I/O operations, e.g. pipes, files and sockets
 - Can be selected when ready to perform I/O operation
- *Selectors* and *selection keys*
 - together with selectable channels define a multiplexed, non-blocking I/O facility. Used to select channels ready for I/O

NIO Packages

<code>java.nio</code>	Buffers, which are used throughout the NIO APIs.
<code>java.nio.channels</code>	Channels and selectors.
<code>java.nio.charset</code>	Character encodings.
<code>java.nio.channels.spi</code>	Service-provider classes for channels.
<code>java.nio.charset.spi</code>	Service-provider classes for charsets.
<code>java.util.regex</code>	Pattern matching using regular expressions.

Buffers

- ***Buffer*** is a container for a fixed amount of data of a specific primitive type; Used by channels
 - Content, data
 - Capacity, size of buffer; set when the buffer is created; cannot be changed
 - Limit, the index of the first element that should not be read or written; $\text{limit} \leq \text{capacity}$
 - Position, the index of the next element to be read or written
 - Mark, the index to which its position will be reset when the reset method is invoked
 - Buffer invariant: $0 \leq \text{mark} \leq \text{position} \leq \text{limit} \leq \text{capacity}$



Some Buffer methods

- **allocateDirect()** Allocates a new direct byte buffer. With direct ByteBuffer, JVM avoid intermediate buffering when performing native I/O operations directly upon the direct buffer.
- **allocate()** Allocate a new buffer.
- **clear()** Clear the buffer, i.e. prepare the buffer for writing data by channel-reads or relative puts (limit = capacity; position = 0)
- **flip()** Prepare the buffer for reading data by channel-writes or relative gets (limit = position; position = 0)
- **rewind()** Prepare the buffer for re-reading data (position = 0)
- **mark()** Set this buffer's mark equal to its position (mark = position)
- **reset()** Reset this buffer's position equal to its mark (position = mark)
- **wrap()** Wrap a given array in a buffer
- **get(), put()** Absolute (index-based) and relative (position-based) get/put data from/into the buffer
- **hasRemaining()** Check whether there are any elements between the current position and the limit

Filling/Draining Buffers

- Filling using wrap or put

```
String s = "Some String";
CharBuffer buf1 = CharBuffer.wrap(s);
CharBuffer buf2 = CharBuffer.allocate(s.length());
// put reversed s in to buf2
for (int i = s.length() - 1; i >= 0; i--) {
    buf2.put(s.charAt(i)); // relative put
} // position in buf2 should be 11 after the loop
```

- Draining using get

```
buf2.flip(); // limit = position; position = 0
String r = "";
while (buf2.hasRemaining())
    r += buf2.get();
}
```


Channels

- ***Channels*** represent connections to various I/O sources, such as pipes, sockets, files, datagrams;
 - operate with buffers and I/O sources: move (read/write) data blocks into / out of buffers from / to the I/O sources;
 - can be blocking/non-blocking, enable ***non-blocking I/O operations***

Some Channel Classes

- **For TCP connections**
 - **SocketChannel**
 - **ServerSocketChannel**
- **For UDP communication**
 - **DatagramChannel**
- **For file access**
 - **FileChannel**

FileChannel

- `java.nio.channels.FileChannel`
 - A channel for reading, writing, mapping, and manipulating a file.
- Can be mapped to a buffer in the main memory
 - `MappedByteBuffer()`
- Has a current position within its file which can be both queried and modified.

Some methods of FileChannel

<code>read (dst, pos)</code> <code>write (src, pos)</code>	Read or write at an absolute position in a file without affecting the channel's position.
<code>map()</code>	Map a region of a file directly into memory.
<code>force()</code>	Force out file updates to the underlying storage device, in order to ensure that data are not lost in the event of a system crash.
<code>transferTo()</code> <code>transferFrom()</code>	Bytes can be transferred from a file to some other channel, and vice versa, in a way that can be optimized by many OSs into a very fast transfer directly to or from the file system cache.

FileChannel Example

```
public class FileChannelTest {  
    public static void main(String[] args) {  
        String filename = "test.txt";  
        try {  
            FileInputStream inf = new FileInputStream(filename);  
            try (FileChannel channel = inf.getChannel()) {  
                MappedByteBuffer buffer = channel.map(FileChannel.MapMode.READ_ONLY,  
                                                       0, channel.size());  
  
                WritableByteChannel out = Channels.newChannel(System.out);  
                while (buffer.hasRemaining()) {  
                    out.write(buffer);  
                }  
            }  
        } catch (IOException e) {  
            e.printStackTrace();  
        }  
    }  
}
```

SocketChannel

- A **selectable channel** for **TCP sockets**.
 - Reads from and writes to a TCP socket.
- Each **SocketChannel** is associated with a **Socket** object

```
SocketChannel channel = SocketChannel.open();  
channel.configureBlocking(false); Non blocking  
channel.connect(new InetSocketAddress(host,  
port));
```

Socket Channel Example

```
public class HTTPClient {
    public static final String GET_REQUEST = "GET / HTTP/1.1\n";

    public static void main(String[] args) {
        String host = (args.length > 0) ? args[0] : "www.kth.se";
        String hostHeader = "Host: " + host + "\n\n";
        int port = (args.length > 1) ? Integer.parseInt(args[1]) : 80;
        WritableByteChannel out = Channels.newChannel(System.out);
        try {
            SocketChannel channel = SocketChannel.open(new InetSocketAddress(
                host, port));
            ByteBuffer buf = ByteBuffer.wrap(GET_REQUEST.getBytes());
            channel.write(buf);
            buf = ByteBuffer.wrap(hostHeader.getBytes());
            channel.write(buf);
            buf = ByteBuffer.allocate(1024);
            while (buf.hasRemaining() && channel.read(buf) != -1) {
                buf.flip();
                out.write(buf);
                buf.clear();
            }
        } catch (IOException e) {
            e.printStackTrace();
            System.exit(0);
        }
    }
}
```

ServerSocketChannel

- *A selectable channel for TCP listening sockets.*
- Each `ServerSocketChannel` is associated with a `ServerSocket` object

```
ServerSocketChannel serverChannel = ServerSocketChannel.open();  
ServerSocket socket = serverChannel.socket();  
socket.bind(new InetSocketAddress(port));
```


Selectors

- ***Selector*** is an object used to select a channel ready to communicate (to perform an operation)
 - Used to operate with several non-blocking channels
 - Allows readiness selection
 - Ability to choose a selectable channel that is ready for some of network operation, e.g. accept, write, read, connect

Selectable Channels

- *Selectable channels* include:
 - DatagramChannel
 - Pipe.SinkChannel
 - Pipe.SourceChannel
 - ServerSocketChannel
 - SocketChannel
- Channels are **registered** with a selector for specific operations, e.g. accept, read, write
- Registration is represented by a *selection key*

Selection Keys

- A selector operates with set of selection keys
- *Selection key* is a token representing the registration of a channel with a selector
- The selector maintains three sets of keys
 - *Key set* contains the keys with registered channels;
 - *Selected-key set* contains the keys with channels ready for at least one of the operations;
 - *Cancelled-key set* contains cancelled keys whose channels have not yet been deregistered.
 - The last two sets are sub-sets of the Key set.

Use of Selectors

- **Create a selector**
`Selector selector = Selector.open();`
- **Configure a channel to be non-blocking**
`channel.configureBlocking(false);`
- **Register a channel with the selector for specified operations (accept, connect, read, write)**
`ServerSocketChannel serverChannel =
 ServerSocketChannel.open();
ServerSocket serverSocket = serverChannel.socket();
serverSocket.bind(new InetSocketAddress(port));
serverChannel.configureBlocking(false); non-Blocking
serverChannel.register(selector, SelectionKey.OP_ACCEPT);`
 - **Register as many channels as you have/need**

Selector methods

- **`select()`** blocking select, returns a set of keys whose channels are ready for I/O.
- **`selectNow()`** non-blocking select, returns zero if no channels are ready
- **`selectedKeys()`** returns the selected-key set
- Iterate over the selected-key set and handle the channels ready for different I/O operations, e.g. read, write, accept

SelectionKey

- Upon registration, each of the registered channels is assigned a selection key.

```
SelectionKey clientKey =  
    clientChannel.register(selector, SelectionKey.OP_READ);
```

- Selection key allows attaching of a single arbitrary object to it

- Associate application data (e.g. a buffer) with the key

```
ByteBuffer buffer = ByteBuffer.allocate(1024);  
clientKey.attach(buffer);
```

- Get the channel and attachment from the key

```
SocketChannel clientChannel =  
    (SocketChannel) key.channel();  
ByteBuffer buffer = (ByteBuffer) key.attachment();
```

Non-Blocking Echo Server

```
while (true) {
    selector.select();
    Iterator<SelectionKey> keys = selector.selectedKeys().iterator();

    while (keys.hasNext()) {
        SelectionKey key = keys.next();
        keys.remove();

        if (key.isAcceptable()) { // accept connection.
            ServerSocketChannel server =
                (ServerSocketChannel) key.channel();
            SocketChannel channel = server.accept();
            channel.configureBlocking(false);
            channel.register(selector, SelectionKey.OP_READ,
                ByteBuffer.allocate(1024));

        } else if (key.isReadable()) { // read from a channel.
            SocketChannel channel = (SocketChannel) key.channel();
            ByteBuffer buffer = (ByteBuffer) key.attachment();
            channel.read(buffer);
            key.interestOps(SelectionKey.OP_WRITE);
        }
    }
}
```

Non-Blocking Echo Server, Cont'd

```
    } else if (key.isWritable()) { // write buffer to channel.
        SocketChannel channel = (SocketChannel) key.channel();
        ByteBuffer buffer = (ByteBuffer) key.attachment();
        buffer.flip();
        channel.write(buffer);
        if (buffer.hasRemaining()) {
            buffer.compact();
        } else {
            buffer.clear();
            key.interestOps(SelectionKey.OP_READ);
        }
    }
}
```